

# Synchronization Tools (cont'd)

Course Code: CSC 2209

Course Title: Operating Systems



**Dept. of Computer Science**  
**Faculty of Science and Technology**

<b>Lecturer No:</b>	<b>09</b>	<b>Week No:</b>	<b>09</b>	<b>Semester:</b>	
<b>Lecturer:</b>	<i>Name &amp; email</i>				

# Lecture Outline



1. Hardware Support for Synchronization
2. Mutex Locks
3. Semaphores

# Synchronization Hardware

- ❑ Many systems provide **hardware support for implementing the critical section (CS) code.**
- ❑ Uniprocessors – could disable interrupts
  - ❑ Currently running code would execute without preemption
  - ❑ Generally **too inefficient on multiprocessor systems**
    - ❑ Operating systems using this not broadly scalable
- ❑ We will look **at three forms of hardware support:**
  1. Memory barriers
  2. Hardware instructions
  3. Atomic variables

# Memory Barriers

- ❑ **Memory models** are the memory guarantees a computer architecture makes to application programs.
- ❑ Memory models may be either:
  - Strongly ordered** – where a memory modification of one processor is immediately visible to all other processors.
  - Weakly ordered** – where a memory modification of one processor may not be immediately visible to all other processors.
- ❑ A **memory barrier** is an instruction that forces any change in memory to be propagated (made visible) to all other processors.

# Memory Barrier

- ❑ We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100:

- ❑ Thread 1 now performs

```
while (!flag)
    memory_barrier();
print x
```

- ❑ Thread 2 now performs

```
x = 100;
memory_barrier();
flag = true
```

# Hardware Instructions

- ❑ Special hardware instructions that allow us to either *test-and-modify* the content of a word, or to *swap* the contents of two words *atomically* (uninterruptibly.)
- ❑ **Test-and-Set** instruction
- ❑ **Compare-and-Swap** instruction

# test\_and\_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    1    boolean rv = *target;    //
    2    *target = true; //
    3    return rv;    //
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to **true**

# Solution using test\_and\_set()

- ❑ Shared boolean variable **lock**, initialized to **false**
- ❑ Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
    /* remainder section */  
} while (true);
```



# compare\_and\_swap Instruction

## Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter **value**
3. Set the variable **value** the value of the passed parameter **new\_value** but only if **\*value == expected** is true. That is, the swap takes place only under this condition.

# Solution using compare\_and\_swap

- ❑ Shared integer **lock** initialized to 0;
- ❑ Solution:

```
while (true){  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

# Bounded-waiting Mutual Exclusion with compare-and-swap

```
while (true) {
    waiting[i] = true;
    key = 1;
    while (waiting[i] && key == 1)
        key = compare_and_swap(&lock, 0, 1);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = 0;
    else
        waiting[j] = false;
    /* remainder section */
}
```

# Atomic Variables

- ❑ Typically, instructions such as **compare-and-swap** are used as building blocks for other synchronization tools.
- ❑ One tool is an **atomic variable** that provides *atomic* (uninterruptible) updates on basic data types such as integers and booleans.
- ❑ For example, the **increment()** operation on the atomic variable **sequence** ensures **sequence** is incremented without interruption:

**increment(&sequence);**

# Atomic Variables

- ❑ The `increment()` function can be implemented as follows:

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp !=
(compare_and_swap(v, temp, temp+1)) );
}
```

# Mutex Locks

- ❑ Previous solutions are complicated and generally inaccessible to application programmers
- ❑ OS designers build software tools to solve critical section (CS) problem
- ❑ Simplest is **mutex lock**
- ❑ Protect a critical section by first **acquire()** a lock then **release()** the lock
  - ❑ Boolean variable indicating if lock is available or not
- ❑ Calls to **acquire()** and **release()** must be atomic
  - ❑ Usually implemented via hardware atomic instructions such as compare-and-swap.
- ❑ But this solution requires **busy waiting- waste CPU cycle**
  - ❑ This lock therefore called a **spinlock**

# Solution to Critical-section Problem Using Locks

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

# Mutex Lock Definitions

```
❑    acquire() {  
        while (!available)  
  
        ; /* busy wait */  
  
        available = false;;  
  
    }  
critical section
```

```
❑    release() {  
  
        available = true;  
  
    }
```

These two functions must be implemented atomically. Both test-and-set and compare-and-swap can be used to implement these functions.



# Semaphore

- ❑ Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- ❑ Semaphore  $S$  – integer variable
- ❑ Can only be accessed via two indivisible (atomic) operations
  - ❑ **wait()** and **signal()**
    - ❑ (Originally called **P()** and **V()**)
- ❑ Definition of the **wait() operation**  
**wait(S) {**  
    **while (S <= 0)**  
        **; // busy wait**  
    **S--;**  
**}**
- ❑ Definition of the **signal() operation**  
**signal(S) { S++; }**

# Semaphore Usage

- ❑ **Counting semaphore** – integer value can range over an unrestricted domain (+infinity to - infinity)
- ❑ **Binary semaphore** – integer value can range only between 0 and 1
  - ❑ Same as a **mutex lock**
- ❑ Can solve various synchronization problems
- ❑ Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$

Create a semaphore “**synch**” initialized to 0

**P1:**

$S_1$ ;

**signal(synch);**

**P2:**

**wait(synch);**

$S_2$ ;

- ❑ Can implement a counting semaphore  $S$  as a binary semaphore

# Semaphore Implementation

- ❑ Must **guarantee that no two processes can execute the wait() and signal() on the same semaphore at the same time**
- ❑ Thus, the implementation becomes the critical section problem where the wait and signal code are placed in the critical section
  - ❑ Could now have **busy waiting** in critical section implementation
    - ❑ But implementation code is short
    - ❑ Little busy waiting if critical section rarely occupied
- ❑ Note that applications may spend lots of time in critical sections and therefore this is not a good solution

# Semaphore Implementation with no Busy waiting

- ❑ With each semaphore there is an associated **waiting queue**
- ❑ Each entry in a **waiting queue** has two data items:
  - ❑ value (of type integer)
  - ❑ pointer to next record in the list
- ❑ Two operations:
  - ❑ **block** – place the process invoking the operation on the appropriate waiting queue
  - ❑ **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ❑ **typedef struct {**  
    **int value;**  
    **struct process \*list;**  
**} semaphore;**

# Implementation with no Busy waiting (cont'd)

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

# Problems with Semaphores

- ❑ Incorrect use of semaphore operations:
  - ❑ `signal (mutex) .... wait (mutex)`
  - ❑ `wait (mutex) ... wait (mutex)`
  - ❑ Omitting of `wait (mutex)` and/or `signal (mutex)`
- ❑ These – and others – are examples of what can occur when semaphores and other synchronization tools are used incorrectly.



# Books

- ❑ Operating Systems Concept
  - ❑ Written by Galvin and Silberschatz
  - ❑ Edition: 9<sup>th</sup>



# References

- ❑ Operating Systems Concept
  - ❑ Written by Galvin and Silberschatz
  - ❑ Edition: 9<sup>th</sup>